

包云泰

(+86) 156-0082-5199 · baoyuntai@outlook.com · GitHub colored-dye

🎓 教育背景

浙江大学(软件学院), 人工智能, 在读博士研究生	2023.09–2028.06
2024–2025 学年获院设优秀实践奖学金。	
浙江大学(计算机科学与技术学院), 信息安全, 工学学士	2019.09–2023.06

▣ 学术成果及经历

- [ICLR 2026] Faithful Bi-Directional Model Steering via Distribution Matching and Distributed Interchange Interventions
Yuntai Bao, Xuhong Zhang, Jintao Chen*, Ge Su, Yuxiang Cai, Hao Peng, Bing Sun, Haiqin Weng, Liu Yan, Jianwei Yin
- [IJCAI 2025] Scalable Multi-Stage Influence Function for Large Language Models via Eigenvalue-Corrected Kronecker-Factored Parameterization
Yuntai Bao, Xuhong Zhang*, Tianyu Du*, Xinkui Zhao, Jiang Zong, Hao Peng, Jianwei Yin
- [Findings of ACL 2025] Probing the Geometry of Truth: Consistency and Generalization of Truth Directions in LLMs Across Logical Transformations and Question Answering Tasks
Yuntai Bao, Xuhong Zhang*, Tianyu Du, Xinkui Zhao, Zhengwen Feng, Hao Peng, Jianwei Yin
- [ACM MM 2022] SongDriver: Real-time Music Accompaniment Generation without Logical Latency nor Exposure Bias
第十作者(共 11 人), 负责评估指标的设计和实验数据预处理。
- [学术服务] 曾在 NeurIPS 2025 Mechanistic Interpretability Workshop 担任志愿审稿人。

🔧 实习经历

浙江大学嘉兴研究院	2022.11–2023.06
本科毕业设计 赵永望教授祥云高安全无人机团队成员	

- 基于 seL4 微内核实现无人机加密通信。本人负责机载任务计算机端及地面控制站端加密通信系统的实现, 包括密钥交换算法、AES 加解密算法及驱动。

▣ 社会实践/其他

- 2021 年 7-8 月参与了计算机学院与灵隐街道联合举办的大学生挂职锻炼社会实践活动, 担任黄龙社区书记助理, 协助社区治理日常事务, 包括慰问空巢老人、消防巡查等。

⚙️ IT 技能

- 编程语言: Python, C/C++, Go, Rust
- 机器学习库: PyTorch, Transformers, sklearn

👤 个人陈述

本人主要研究领域为 LLM 机制可解释性 (mechanistic interpretability), 尤其是借助可解释性的技术和见解实现有效 (effective)、高效 (efficient) 且鲁棒 (robust) 的模型控制, 当下主要研究基于表征控制的模型引导 (representation steering)。作为以可解释性见长的研究者, 我拥有一般 LLM 研究者不具备的**比较优势**: 我对 LLM 的内部工作机制有较深入的技术积累和见解, 这些经验可以进而帮助我提出比传统机器学习路线更有针对性 (targeted)、更精细 (granular)、数据效率 (data efficiency) 更高的解决方案。

REFERENCES

- Yuntai Bao, Xuhong Zhang, Tianyu Du, Xinkui Zhao, Zhengwen Feng, Hao Peng, and Jianwei Yin. Probing the geometry of truth: Consistency and generalization of truth directions in LLMs across logical transformations and question answering tasks. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Findings of the Association for Computational Linguistics: ACL 2025*, pp. 682–700, Vienna, Austria, July 2025a. Association for Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.38. URL <https://aclanthology.org/2025.findings-acl.38/>.
- Yuntai Bao, Xuhong Zhang, Tianyu Du, Xinkui Zhao, Jiang Zong, Hao Peng, and Jianwei Yin. Scalable multi-stage influence function for large language models via eigenvalue-corrected kronecker-factored parameterization. In James Kwok (ed.), *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence, IJCAI-25*, pp. 8022–8030. International Joint Conferences on Artificial Intelligence Organization, 8 2025b. doi: 10.24963/ijcai.2025/892. URL <https://doi.org/10.24963/ijcai.2025/892>. Main Track.
- Yuntai Bao, Xuhong Zhang, Jintao Chen, Ge Su, Yuxiang Cai, Hao Peng, Bing Sun, Haiqin Weng, Liu Yan, and Jianwei Yin. Faithful bi-directional model steering via distribution matching and distributed interchange interventions. *arXiv preprint arXiv:2602.05234*, 2026.
- Zihao Wang, Kejun Zhang, Yuxing Wang, Chen Zhang, Qihao Liang, Pengfei Yu, Yongsheng Feng, Wenbo Liu, Yikai Wang, Yuntao Bao, and Yiheng Yang. Songdriver: Real-time music accompaniment generation without logical latency nor exposure bias. In *Proceedings of the 30th ACM International Conference on Multimedia, MM ’22*, pp. 1057–1067, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450392037. doi: 10.1145/3503161.3548368. URL <https://doi.org/10.1145/3503161.3548368>.